

Phloem MPI Benchmarks

Summary Version

1.4.1

Purpose of Benchmark

The purpose of the Phloem MPI benchmark suite is to benchmark the bandwidth, latency, and messaging rate of basic communication operations.

Characteristics of Benchmark

The Phloem MPI benchmark suite consists of three independent benchmarks: Presta, mpiBench, and SQMR.

The Presta benchmark measures ping-pong latency and aggregate bandwidth for 1 or more pairs of MPI processes to provide intra- and inter-node aggregate bandwidth as well as bisection bandwidth.

The mpiBench benchmark measures collective latency for several blocking and non-blocking collectives for MPI_COM_WORLD and sub-communicators.

The SQMR benchmark measures messaging rate for MPI point-to-point operations.

Mechanics of Building Benchmark

Build flags can be modified in the top-level Makefile.inc file. The default make target should build the appropriate benchmarks with “make”.

Mechanics of Running Benchmark

Please see the Phloem top-level README as well as individual benchmark READMEs for more information regarding the benchmarks. The following items identify measurements of interest and example commands for running the benchmarks.

MPI point-to-point intra-node aggregate bandwidth

Run the presta/com benchmark with the number of cores MPI processes on a single node.

```
mpirun -n 16 ./com -m bw.message.sizes # on 1 node
```

MPI point-to-point inter-node aggregate bandwidth

Run the presta/com benchmark with MPI processes equal to 2x the number of cores per node over two nodes.

```
mpirun -n 32 ./com -m bw.message.sizes # over 2 nodes
```

MPI bi-section bandwidth

Run the presta/com benchmark with MPI processes equal to Nx the number of cores per node over all nodes, where N is the number of nodes in the system.

```
mpirun -n N ./com -m bw.message.sizes # over all nodes
```

MPI point-to-point intra-node latency

Run the presta/com benchmark with the number of cores MPI processes on a single node.

```
mpirun -n 16 ./com -m latency.message.sizes -w Latency # on 1 node
```

MPI point-to-point inter-node latency

Run the presta/com benchmark with MPI processes equal to 2x the number of cores per node over two nodes.

```
mpirun -n 32 ./com -m latency.message.sizes -w Latency # over 2 nodes
```

MPI Collective Latency

Run the mpigraph/mpiBench benchmark with MPI processes equal to Nx the number of cores per node over all nodes, where N is the number of nodes in the system.

```
mpirun -n 32 ./mpiBench -d 2 -p 2
```

MPI Messaging Rate

Measure the messaging rate for a single MPI process receiving messages from a single remote MPI processes. This should be run with 1 process per node over 2 nodes.

```
mpirun -n 2 ./sqmr --num_cores=1 --num_nbors=1
```

Measure the messaging rate for a single MPI process receiving messages from multiple MPI processes. This should be run with 1 process per node over 4 nodes.

```
mpirun -n 4 ./sqmr --num_cores=1 --num_nbors=3
```

Measure the aggregate messaging rate for MPI processes on a node receiving messages from multiple MPI processes. This should be run with multiple processes per node over multiple nodes. Example for 4 processes per node over 4 nodes:

```
mpirun -n 16 ./sqmr --num_cores=4 --num_nbors=3
```